

# **Design, Develop, and Deploy**

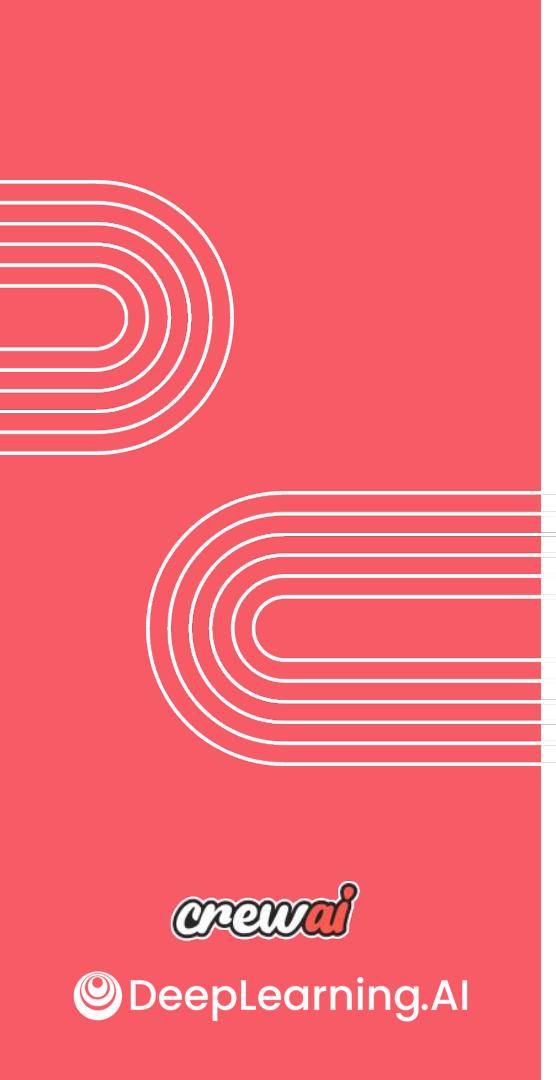
## **Multi-Agent Systems**

### **with CrewAI**

---

Managing Systems of AI Agents





# Managing Systems of AI Agents

---

## Collaboration


*crewai*

 DeepLearning.AI


# Single Agent System




# Single Agent System




# Single Agent System



# Multi-Agent System



# Breaking Down Complexity



# Agent

Role

Real-world job title


Goal

What “good” looks like

Backstory

Expertise, working style, values

# Example



# Specialists > Generalists

Focused agents deliver precise outputs

**Example:**

**Technical Blog Writer Agent** instead of **Writer Agent**  
If the goal is to have a technical blog creation agent.

- Specialize in role, versatile in application
- Complementary skills across the crew

# Agent

Role

Real-world job title

Goal

What “good” looks like

Backstory

Expertise, working style, values

LLM

Which model to use

## **Smaller models**

- *Cheaper*
- *Faster*

**Speed**

**Quality**

## **Bigger Models**

- *More capable reasoning*
- *Follow instructions better*

### **Smaller models**

- *Cheaper*
- *Faster*

**Speed**

**Quality**

**Consistency**

### **Bigger Models**

- *More capable reasoning*
- *Follow instructions better*

### **Smaller models**

- *Cheaper*
- *Faster*

**Speed**


**Quality**

**Consistency**

### **Bigger Models**

- *More capable reasoning*
- *Follow instructions better*

*Speed*



---

**Smaller models**

**Bigger models**

## **Smaller models**

- *Cheaper*
- *Faster*


**Speed**

**Quality**

**Consistency**


## **Bigger Models**

- *More capable reasoning*
- *Follow instructions better*



**Smaller models**

**Bigger models**



# Managing Systems of AI Agents

---

## Communication

*crewai*


 DeepLearning.AI

# Agents vs Crews

## Single Agent




## Multi-Agent




# Sequential Process

*Crew*



Output of each Agent's Task serves  
as context for the next one.

# Sequential Process




```
# Create a planning task
plan = Task(description="Develop a detailed plan to create a Ruby on Rails app...",
            expected_output="A detailed product specification document...",
            agent=planning_agent)

# Create a testing task
tests = Task(description="Create a set of unit tests for this app...",
              expected_output="A set of runnable unit tests...",
              agent=tester_agent,
              context=[plan])


# Create a app
app = Task(description="Develop a Ruby on Rails app using the specs, that passes the tests...",
           expected_output="A fully functional Ruby on Rails app...",
           agent=coding_agent,
           context=[plan, tests])
```

# Parallel Process



# Hierarchical Process


*Crew*




# Using Memory for Communication between Agents




# Using Memory for Communication between Agents



# Using Memory for Communication between Agents



# Using Memory for Communication between Agents



# Common Coordination Issues

## Redundant work:


Tighten the scope of tasks each agent performs to avoid overlap; start with parallel process then hierarchical when needed.

## Slow serial chains:


Mark independent tasks **async\_execution=True** and gate with a later task's **context**

# Process Types


Sequential




Hierarchical



Hybrid



Parallel






# Managing Systems of AI Agents

---


Building  
Coordination Patterns


*crewai*

# Deep Research Crew (Sequential)



# Deep Research Crew (Hybrid)





# Managing Systems of AI Agents

---

## Using the A2A Protocol

crewai

DeepLearning.AI

# Agent2Agent Protocol

Agent2Agent Protocol (**A2A**) is an open standard for AI agent communication across different frameworks and vendors

## simplicity

Decouple public specifications and private implementations of agents

## standardization

Promotes interoperability between agents without wrapping them as tools

## security

Built on familiar constructs like HTTP / JSON-RPC with support for authentication

## Key components of A2A

**Agent Card**

Document describing agent identity, capabilities, and connection details

**Task**

Stateful unit of work supporting long horizon jobs

**Message**

Single turn of conversation between agents and/or users

**Artifact**

Output from tasks consisting of multimodal data

# A2A versus MCP

- A2A complements MCP
- Use MCP to connect agents to tools
- Use A2A to connect agents to agents
- Wrapping agents as tools could limit their capabilities for collaboration and communication


MCP

Stateless standard for connecting agent to tool, prompts or datasets


A2A


Stateful standard for interactions between different agents

# Adopting A2A



# Adopting A2A





# Managing Systems of AI Agents

---

## Orchestrating Agents with Flows

*crewai*

# Mental Models for Agentic Systems

## Agents

Real-time planning


## Graphs

Nodes and edges


## Events

Trigger-based workflows

# Flows




# Flows




## Crew

*more agency*



## Flows


*finer precision*



# Opt-In Agency Model

What is the minimum I  
can get away with for my  
use-case?

- Structured backbone
- Add agency according to the need
- Start with LLM and expand to full Crew



# Building Blocks of Flows

## @start

First function that is executed

## @listen

Makes a method execute when the specified task generates an output


## @router


Defines conditional routing logic based on an input


```
from crewai.flow.flow import listen, start
```


```
@start()  
def start_conversation(self):  
    print("Starting conversation")
```


```
@listen(start_conversation)  
def trigger_research(self):  
    print("Checking for research")
```











C:>  
crewai flow plot

Get User Message

Check For Deep Research

Answer

Clarify

Start Research

Generate Report

Answer

Clarify

Start Research

Generate Report

C:>

```
crewai create flow deep_research
```

```
+--- pyproject.toml  
+--- tests  
+--- README.md  
+--- .gitignore  
+--- .env  
+--- src  
    +--- deep_research  
        +--- crews  
            |   +--- research_crew  
            |   |   +--- research_crew.py  
            |   |   +--- config  
            |   |   |   +--- agents.yaml  
            |   |   |   +--- tasks.yaml  
            |   |   +--- __init__.py  
        +--- tools  
            |   +--- __init__.py  
            |   +--- custom_tool.py  
        +--- __init__.py  
    +--- main.py
```

# Adding State to Flows

State provides **shared context** across each step of your flow

## During Execution

- Each step of the flow consists of a function with access to state
- All functions can read / write state throughout execution
- Accumulated data in state can inform routing of the flow

## After Execution

- Optionally you can persist state. Persistence store state for later use.
- Note persistent state of flows is different from the memory of crews in the flow
- Using persistence is especially important with conversational agents

# Making State Persistent

Persistent state allows for flows that can be **paused**, **resumed**, and even **recover** after failures

## Class Level Persistence

- When applied at the class level, the state updates and saves after every function execution in the flow

## Method Level Persistence

- When applied at the method level, the state updates and saves after execution of functions with `@persist()` decorator




# Managing Systems of AI Agents


---


Building a  
Deep Research Flow

*crewai*

# Deep Research Crew (Hybrid)







# Managing Systems of AI Agents

---

Tactics for Building  
Reliable Systems

crewai

You can build an agent  
*But can it run reliable every time?*

More importantly: Can you trust it?

# How to Build Agents you Trust

**Including Deterministic  
Controls to Agentic system**


Flows  
Guardrails  
Reasoning agents  
Human-in-the-loop oversight  
Testing  
Training  
Structured output  
Safe code execution

# **Adding reasoning**

***Requiring your agents to pre-plan***

**Turn on Reasoning at the Agent level  
to force a Pre-Plan before execution**

# Reasoning




```
from crewai import Agent

# Create an analyst agent with reasoning enabled
analyst = Agent(
    role="Data Analyst",
    goal="Find actionable insights from the data",
    backstory="Experienced at EDA and communication",
    reasoning=True,
    max_reasoning_attempts=3,
    verbose=True
)
```

# **Adding Human in the Loop**

***Strong code check for tasks outputs***

# Human in the Loop



# Human in the Loop

C:>

=====

## HUMAN FEEDBACK: Provide feedback on the Final Result and Agent's actions.

Please follow these guidelines:


- If you are happy with the result, simply hit Enter without typing anything.
- Otherwise, provide specific improvement requests.
- You can provide multiple rounds of feedback until satisfied.

=====


# **Training your crew**

***Building iterative feedback into memory***

# Training



# Training



# Training


```
C:>  
crewai train
```

```
DeepResearchFlow().crew().train(  
    n_iterations=n_iterations,  
    inputs=inputs,  
    filename=filename  
)
```

# **Running tests for your crew**

***Comparing performance of different LLMs***

# Testing




# Testing

```
C:>  
crewai test
```

```
DeepResearchFlow().crew().test(  
    n_iterations=n_iterations,  
    model=model_id,  
)
```

# Testing

## Crew



Tasks Scores (1-10 Higher is better)				
Tasks/Crew	Run 1	Run 2	Run 3	Avg. Total
Task 1	7.0	6.5	8.0	7.2
Task 2	5.5	7.0	6.0	6.2
Crew	7.0	5.0	6.5	6.3
Execution Time (s)	45	40	59	48

# **Safe code execution**


***Running code written by agents***

# Code Interpreter Tool

```
from crewai import Agent

programmer_agent = Agent(
    role="Python Programmer",
    goal="Write and execute Python code to solve problems",
    backstory="An expert Python programmer.",
    allow_code_execution=True,
    code_execution_mode="safe",
)
```

**Reliable agents aren't just accurate  
they're *predictable, measurable, and recoverable***



# Managing Systems of AI Agents

---

## Monitoring and Observability

*crewai*

 DeepLearning.AI

# How to Build Agents you Trust

Debuggability

Quality Monitoring

Data Governance

Protecting PII

Prompt Injection Safeguards

Secure Code Generation

## Observability

Debuggability  
Quality Monitoring

## Security

Data Governance  
Prompt Injection Safeguards  
Secure Code Generation

## Compliance

Protecting PII

**“You can’t fix what you can’t see.”**

# Tracing

```
from crewai import Agent, Crew, Process, Task
from crewai_tools import SerperDevTool

# Define your first agent
researcher = Agent(
    role="AI LLMs Senior Data Researcher",
    goal="Uncover cutting-edge developments in AI and data science",
    backstory="""You work at a leading tech think tank.
    Your expertise lies in identifying emerging trends.
    You have a knack for dissecting complex data and presenting actionable insights.""",
    verbose=True,
    tools=[SerperDevTool()],
)
```

# Tracing

```
# Define your second agent
writer = Agent(
    role="AI LLMs Reporting Analyst",
    goal="Craft compelling content on tech advancements",
    backstory="""You are a renowned Technical Writer, known for your insightful and engaging
articles.

    You transform complex concepts into compelling narratives.""",
    verbose=True,
)
```

# Tracing

```
# Create tasks for your agents

research_task = Task(
    description="""Conduct a comprehensive analysis of the latest advancements in LLMs.
    Identify key trends, breakthrough technologies, and potential industry impacts.""",
    expected_output="Full analysis report in bullet points",
    agent=researcher,
)
reporting_task = Task(
    description="""Using the insights provided, develop an engaging white paper
    that highlights the most significant AI advancements.
    Your post should be informative yet accessible, catering to a tech-savvy audience.""",
    expected_output="white paper with at least 4 paragraphs",
    agent=writer,)
```

# Tracing

```
# Enable tracing in your crew
crew = Crew(
    agents=[researcher, writer],
    tasks=[research_task, writing_task],
    process=Process.sequential,
    tracing=True, # Enable built-in tracing
    verbose=True
)

# Execute your crew
result = crew.kickoff()
```

▼	👤 AI LLMs Senior Data Researcher	
	0.01s (+9.43s) • 1 task	
▼	📝 research_task	⌚ 9.43s +0.01s
	📄 Started	+0.00s
	🕒 LLM call	⌚ 9.4s +9.42s
	📋 Completed	+9.43s
▼	👤 AI LLMs Reporting Analyst	
	9.44s (+21.51s) • 1 task	
▼	📝 reporting_task	⌚ 21.51s +9.44s
	📄 Started	+0.00s
	🕒 LLM call	⌚ 21.5s +21.49s
	📋 Completed	+21.51s

## ⌚ Timeline

▼ AI LLMs Senior Data Researcher  
0.01s (+9.43s) • 1 task

▼ research\_task

Started

⌚ 9.43s +0.01s

+0.00s

LLM call

⌚ 9.4s +9.42s

Completed

+9.43s

▼ AI LLMs Reporting Analyst  
9.44s (+21.51s) • 1 task

▼ reporting\_task

Started

⌚ 21.51s +9.44s

+0.00s

LLM call

⌚ 21.5s +21.49s

Completed

+21.51s

## ⌚ Event details

Details Messages Raw Data

### LLM call

9/9/2025, 12:34:14 AM

### Response

I now can give a great answer

Final Answer:

1. **Model Size and Efficiency**: In 2025, several AI LLMs have exceeded 1 trillion parameters, with optimized architectures that improve computational efficiency. Researchers are focusing on pruning and quantization techniques to deploy these large models on edge devices without significant performance loss.

2. **Multimodal Capabilities**: Leading LLMs can now seamlessly integrate text, images, and audio, leading to advanced applications in fields like virtual reality, education, and creative arts. These models demonstrate a better understanding of contextual relationships across different modalities.

3. **Personalization Features**: AI LLMs are now equipped with adaptive learning capabilities that allow them to personalize interactions based on user behaviors and preferences, resulting in more engaging and user-centered experiences in applications like virtual assistants and customer service.

4. **Interdisciplinary Applications**: AI LLMs are being applied across various fields, including healthcare for diagnostic support, law for contract analysis, and climate science for predictive modeling, showcasing their versatility and enhanced role in decision-making processes.

## Timeline

▼ AI LLMs Senior Data Researcher  
0.00s (+16.30s) • 1 task

### ▼ research\_task

Started

⌚ 16.30s +0.00s

+0.00s

LLM call

⌚ 1.4s +1.42s

Search the internet with Serper

+1.43s

LLM call

⌚ 999ms +4.12s

Read website content

+4.12s

LLM call

⌚ 1.7s +6.48s

Read website content

+6.49s

LLM call

⌚ 7.6s +16.29s

+16.30s

Completed

## Event details

Details Raw Data

### ⌚ Search the internet with Serper

9/9/2025, 12:44:17 AM

#### ▼ Input

```
{"search_query": "latest developments in AI LLMs 2025"}
```

Running

Started 12:44:17 AM

# **Production Agents Operate at Machine Speed**

*So observability becomes reactive  
and proactive alerting becomes larger priority*

---

## Crew Quality Sampling

---

Sampling 10% of executions for quality checks.

### Sampling Percentage

Select the portion of executions to be evaluated.



### Enable Quality Sampling

Toggle this feature to enable or disable quality sampling for this deployment.

Enabled

### ▼ Advanced thresholds

Performance Threshold (0-10)


7


Max Execution Time (seconds)

3600

## Tokens, Execution Time and Test Results

Aggregating prompt tokens and execution time. Displaying averaged daily test results.





# Managing Systems of AI Agents

---

CI/CD for Agents

crewai

 DeepLearning.AI

# CI & CD for Agentic Systems

- How can you support new models?
- How do these systems evolve over time?
- How do you ensure they are still working?

# CI & CD for Agentic Systems



# Prompts

(Roles, Goals, Backstories, Descriptions  
and Expected Outputs)

# Agent Logic

(Custom tools, Hooks, Guardrails  
and such)

```
C:>  
crewai create crew research_crew
```

```
C:>  
cd research_crew
```

```
research_crew/  
└── .gitignore  
└── pyproject.toml  
└── README.md  
└── .env  
└── src/  
    └── research_crew/  
        ├── __init__.py  
        ├── main.py  
        ├── crew.py  
        └── tools/  
            ├── custom_tool.py  
            └── __init__.py  
└── config/  
    ├── agents.yaml  
    └── tasks.yaml
```

# Agent Repository

 **Agents Repository**  
Manage your agents - Create and configure agents to automate tasks and workflows 

 **Agent Repository**

1. Create agents with specific roles and goals for your workflows
2. Configure tools and capabilities for each specialized assistant
3. Deploy agents across projects via visual interface or API integration

 **Your Pre Made Agents**

ROLE	GOAL	TOOLS	ACTIONS
Sales Research	gather sales data about a lead	crewai-tools: ScrapeWebsiteTool   crewai-tools: SerperDevTool   crewai-tools: CrewaiEnterpriseTools	 

# Agent Repository

```
from crewai import Agent  
agent = Agent(from_repository="sales-research")
```

# Tool Repository



## Tools

Collaborate by sharing tools within your organization, or publish them publicly to contribute with the community.

### + Getting Started

1. Creating a new tool? Use `crewai tool create your-tool`
2. Ready to share your tool? Publish with `crewai tool publish`
3. Want to use a tool? Install it with `crewai tool install your-tool`



## Your Tools

### crewai-tools Public

CrewAI Tools - Official collection of tools for CrewAI

### your-tool Private

Power up your crews with your\_tool

# Don't Overlook the Planning

1. Not spending time on planning use cases
2. Not clear definition of success
3. Not breaking the process into smaller chunks
4. Not measuring / evaluating